




Asko Soukka, asko.soukka@iki.fi
github.com/datakurre

ZODB

the Python database



JYVÄSKYLÄN YLIOPISTO



ZODB

the Python database

PyCon Finland 31.10.2016

(Zope) Object Database

- Native object tree database for Python since 2002
- ACID transactions with snapshot isolation (MVCC)
- ZODB, persistent, BTrees, transaction, zope.interface, ZConfig, zc.lockfile, zodbpickle (optional: ZEO, zc.zrs)
- Latest version: 5.0.0 (2016-09-06), Python 2.7, 3.4, 3.5

Persisting TaskList with ZODB

Creating new ZODB

```
In [5]: import ZODB
import ZODB.FileStorage
import transaction

ZODB_FILENAME = 'mydata.fs'

storage = ZODB.FileStorage.FileStorage(ZODB_FILENAME)
db = ZODB.DB(storage)
connection = db.open()
root = connection.root()
```

Analyzing empty ZODB

```
In [6]: from ZODB.scripts import analyze
analyze.report(analyze.analyze(ZODB_FILENAME),)
```

Processed 1 records in 1 transactions
 Average record size is 68.00 bytes
 Average transaction size is 68.00 bytes
 Types used:

Class Name	Count	TBytes	Pct	AvgSize
-----	-----	-----	-----	-----
persistent.mapping.PersistentMapping	1	68	100.0%	68.00
-----	-----	-----	-----	-----
Total Transactions	1			0.07k
Total Records	1	0k	100.0%	68.00
Current Objects	1	0k	100.0%	68.00

Persisting TaskList

```
In [7]: root['tasks'] = TaskList()

import transaction
transaction.commit()
```

Crash Course

ZODB can be a lot of fun


- It's implicit — if object pickles, it persists.

Just remember to inherit your classes from Persistent.

- It's pythonic and does object references right.
- It's safe, append only, immutable, but packs on request.
- It's fast, unless your code slows it down...

Not necessarily *web scale*

- Only one simultaneous request per connection.
 - Each connection has its own *connection cache*.
 - *Connection cache* = unpickled objects in RAM.
 - Must be big enough to store the current working set.
- Search indexes in ZODB are prone to write conflicts.
- No multi-master replication. But see RelStorage and neoppod.



BUILT-IN SECURITY

Make highly granular security assertions against database content, allowing for tight access control of specific resources.

TEXT INDEXING

Built-in text indexing and search allows your users to find their content without requiring you to configure, manage and learn an external text indexing system.

WORKFLOW

Define workflow states and transitions for types of content, and enforce them.



UNDO

Undo virtually any action taken via the management interface or by a retail user.



Substance D


Performance and availability monitoring is baked in. Use any *statsd* provider to set up visibility into your application's performance.



AUDITING

Every action taken by each site user is logged into a built-in audit log facility.

SCREENSHOTS





Questions?

github.com/datakurre/pyconfi2016